

The IP addresses for these systems are set statically. Do not change these settings.

Hostname resolution is configured to resolve the fully qualified hostnames listed above, but also to resolve host shortnames.

account information

The root password for all systems is **flectrag**.

Do not change the root password. Unless otherwise specified, this will be the password used to access other systems and services. Also, unless otherwise specified, this password also applies to all accounts you create, or any service that requires a password to be set.

For convenience, SSH keys are preinstalled on all systems, allowing root access via SSH without entering a password. Do not make any modifications to the root SSH configuration file on the system.

The user account greg has been created on the Ansible control node. This account is pre-installed with SSH keys that allow SSH logins between the Ansible control node and each Ansible managed node. Do not make any modifications to the greg SSH configuration file on the system. You can use su from the root account to access this user account.

Important information

Unless otherwise specified, all your work (including Ansible playbooks, configuration files, host inventory, etc.) should be kept in the directory **/home/greg/ansible** on the control node and should be owned by the greg user. All Ansible related commands should be run by user **greg** from this directory on the Ansible control node.

Additional Information

Some exam items may require modification of the Ansible host inventory. It is your responsibility to ensure that all previous checklist groups and projects are preserved, coexisting with any other changes. You will also have to make sure that all default groups and hosts in the inventory retain any changes you make.

The firewall on the exam system is not enabled by default, and SELinux is in enforcing mode.

If additional software needs to be installed, your physical system and Ansible control node may already be set up to point to the following repositories on content:

http://content/rhel8.0/x86_64/dvd/BaseOS

http://content/rhel8.0/x86_64/dvd/AppStream

Some projects require additional files, which are already available at:

<http://materials>

Product documentation can be found at:

<http://materials/docs/ansible/html>

Other resources are also configured for you to use during the exam. Specific information about these resources will be provided in projects that require them.

Important information

Note that before scoring, your Ansible managed node system will be reset to the initial state at the start of the exam, and the Ansible playbook you have written will be run from the `/home/greg/ansible` directory on the control node by running as user `greg` to apply. After the playbook is running, your managed nodes are evaluated to see if they are configured as specified.

 [RHCE exam questions](#) 

Do all the following steps on your system.

Install and configure Ansible

Create and run Ansible ad hoc commands

Install the package

Using RHEL system roles

Install roles using Ansible Galaxy

Create and use roles

Using roles from Ansible Galaxy

Create and use logical volumes

Generate hosts file

Modify file content

Create a web content directory

Generate hardware report

Create a password vault

Create user account

Update keys for Ansible repositories

😊 Install and configure Ansible

Install and configure Ansible on the control node **172.25.250.254** as described below:

Install the required packages

Create a static inventory file named **/home/greg/ansible/inventory** to meet the following requirements:

172.25.250.9 is a member of the **dev** host group

172.25.250.10 is a member of the **test** host group

172.25.250.11 and **172.25.250.12** are members of the **prod** host group

172.25.250.13 is a member of the **balancers** host group

The **prod** group is a member of the **webservers** host group

Create a configuration file named **/home/greg/ansible/ansible.cfg** to meet the following requirements:

The host inventory file is **/home/greg/ansible/inventory**

The locations of roles used in playbooks include **/home/greg/ansible/roles**

Problem solving method:

```
[greg@bastion ~]$ sudo yum install -y ansible
[greg@bastion ~]$ mkdir -p /home/greg/ansible
[greg@bastion ~]$ cd /home/greg/ansible
[greg@bastion ansible]$ vim inventory
[dev]
172.25.250.9
[test]
172.25.250.10
[prod]
172.25.250.11
172.25.250.12
[balancers]
172.25.250.13
[webservers:children]
prod
[all:vars]
ansible_user=root
ansible_password=redhat
[greg@bastion ansible]$ cp /etc/ansible/ansible.cfg ./
[greg@bastion ansible]$ mkdir roles
[greg@bastion ansible]$ vim /home/greg/ansible/ansible.cfg
#Modify only four lines in total
inventory = /home/greg/ansible/inventory
```

```
roles_path = /home/greg/ansible/roles
host_key_checking=False
remote_user = root
[greg@bastion ansible]$ ansible --version
[greg@bastion ansible]$ ansible-inventory --graph
```

😁 Create and run Ansible ad hoc commands

As a system administrator, you need to install software on managed nodes.

Follow the main text to create a shell script called `/home/greg/ansible/adhoc.sh` that will install the yum repository on each managed node using Ansible ad hoc commands:

Repository 1:

The name of the repository is `EX294_BASE`

Described as `EX294 base software`

The base URL is `http://content/rhel8.0/x86_64/dvd/BaseOS`

GPG signature checking is `enabled`

The GPG key URL is `http://content/rhel8.0/x86_64/dvd/RPM-GPG-KEY-redhat-release`

Repository is `enabled`

Repository 2:

The name of the repository is `EX294_STREAM`

Described as `EX294 stream software`

The base URL is `http://content/rhel8.0/x86_64/dvd/AppStream`

GPG signature checking is `enabled`

The GPG key URL is `http://content/rhel8.0/x86_64/dvd/RPM-GPG-KEY-redhat-release`

Repository is `enabled`

Problem solving method:

```
[greg@bastion ansible]$ ansible-doc yum_repository
[greg@bastion ansible]$ vim adhoc.sh
#!/bin/bash
ansible all -m yum_repository -a 'name="EX294_BASE" description="EX294 base software"
baseurl="http://content/rhel8.0/x86_64/dvd/BaseOS" gpgcheck=yes enabled=1 gpgkey="http:
//content/rhel8.0/x86_64/dvd/RPM-GPG-KEY-redhat-release"'
ansible all -m yum_repository -a 'name="EX294_STREAM" description="EX294 stream software"
baseurl="http://content/rhel8.0/x86_64/dvd/AppStream" gpgcheck=yes enabled=1
gpgkey="http://content/rhel8.0/x86_64/dvd/RPM-GPG-KEY-redhat-release"'
[greg@bastion ansible]$ chmod +x adhoc.sh
```

🤖 Install the package

Create a playbook called `/home/greg/ansible/packages.yml` :

Install the `php` and `mariadb` packages to hosts in the `dev` , `test` and `prod` host groups

Install the `RPM Development Tools` package group on hosts in the `dev` host group

Update all packages on hosts in the `dev` host group to the latest version

Problem solving method:

```
---
- name: install package A
  hosts: dev,test,prod
  tasks:
    - name: one
      yum:
        name: php
        state: latest
    - name: two
      yum:
        name: mariadb
        state: latest
- name: install package B
  hosts: dev
  tasks:
    - name: one
      yum:
        name: "@RPM Development Tools"
        state: latest
    - name: two
      yum:
        name: "*"
        state: latest
```

🤖 Use RHEL system roles

Install the RHEL system role package and create a playbook `/home/greg/ansible/timesync.yml` with the following conditions :

run on `all managed nodes`

Use the `timesync` role

Configure the role to use the currently valid NTP provider

Configure the role to use the time server `172.25.254.254`

Configure the role to enable the `iburst` parameter

Problem solving method:

```
[greg@bastion ansible]$ sudo yum install rhel-system-roles
[greg@bastion ansible]$ vim ansible.cfg
#Modify only one line
```

```

roles_path = /home/greg/ansible/roles:/usr/share/ansible/roles
[greg@bastion ansible]$ ansible-galaxy list
[greg@bastion ansible]$ cp /usr/share/doc/rhel-system-roles/timesync/example-timesync-
playbook.yml timesync.yml
[greg@bastion ansible]$ vim timesync.yml
---
- hosts: all
  vars:
    timesync_ntp_servers:
      - hostname: 172.25.254.254
        iburst: yes
  roles:
    - rhel-system-roles.timesync

```

Install roles using Ansible Galaxy

Use Ansible Galaxy and the requirements file `/home/greg/ansible/roles/requirements.yml` . Download roles from the following URL and install to `/home/greg/ansible/roles` :

`http://materials/haproxy.tar` This role should be named **balancer**

`http://materials/phpinfo.tar` This role should be named **phpinfo**

Problem solving method:

```

[greg@bastion ansible]$ mkdir roles
[greg@bastion ansible]$ cd roles
[greg@bastion roles]$ vim requirements.yml
---
- src: http://materials/haproxy.tar
  name: balancer
- src: http://materials/phpinfo.tar
  name: phpinfo
[greg@bastion roles]$ ansible-galaxy install -r requirements.yml
[greg@bastion roles]$ ansible-galaxy list

```

Create and use characters

Create a role named apache in `/home/greg/ansible/roles` with the following requirements :

httpd package installed, set to **enable and start at system boot**

Firewall is enabled and running with rules that allow access to the **web** server

The template file `index.html.j2` already exists to create the file `/var/www/html/index.html` with the following output :

Welcome to HOSTNAME on IPADDRESS

where **HOSTNAME** is the fully qualified domain name of the managed node and **IPADDRESS** is the IP address of the managed node.

Lao Liu's warm reminder: The role can be created. There is no need to perform a separate call test. Once it is called and used, it will cause the following question to report an error. The following question will immediately call the apache role, just follow the webpage content and steps~

Problem solving method:

```
[greg@bastion roles]$ ansible-galaxy init apache
[greg@bastion roles]$ vim apache/tasks/main.yml
---
- name: one
  yum:
    name: httpd
    state: latest
- name: two
  service:
    name: httpd
    state: started
    enabled: yes
- name: three
  firewallld:
    service: http
    permanent: yes
    state: enabled
    immediate: yes
- name: four
  template:
    src: index.html.j2
    dest: /var/www/html/index.html
[greg@bastion roles]$ vim apache/templates/index.html.j2
Welcome to {{ ansible_fqdn }} on {{ ansible_default_ipv4.address }}
```

😊 Using roles from Ansible Galaxy

Create a playbook named `/home/greg/ansible/roles.yml` with the following requirements :

The playbook contains a play that runs on hosts in the **balancers** host group and will use the **balancer** role.

This role configures a service to load balance web server requests among hosts in the **webservers** host group.

Browsing to a host in the **balancers** host group (eg `http://172.25.250.13`) will generate the following output:

Welcom to serverb.lab.example.com on 172.25.250.11

Reloading the browser will generate output from another web server:

`Welcom to serverc.lab.example.com on 172.25.250.12`

The playbook contains a play that runs on a host in the **webservers** host group and will use the **phpinfo** role.

Browsing to a host in the **webservers** host group at the URL `/hello.php` will generate the following output:

Hello PHP World from FQDN

where FQDN is the fully qualified name of the host.

Hello PHP World from

serverb.lab.example.com

There are also various details of the PHP configuration, such as the installed PHP version, etc.

Likewise, browsing to <http://172.25.250.12/hello.php> produces the following output:

Hello PHP World from

serverc.lab.example.com

There are also various details of the PHP configuration, such as the installed PHP version, etc.

Problem solving method:

```
[greg@bastion ansible]$ vim roles.yml
---
- name: one
  hosts: balancers
  roles:
    - balancer
- name: two
  hosts: webservers
  roles:
    - phpinfo
- name: three
  hosts: webservers
  roles:
    - apache
```

Create and use logical volumes

Create a playbook named [/home/greg/ansible/lv.yml](#) that will run on all managed nodes to perform the following tasks:

Create logical volumes that meet the following requirements:

Logical volumes are created in the [research](#) volume group

The logical volume name is [data](#)

Logical volume size is [1500 MiB](#)

Format a logical volume with the [ext4](#) filesystem

An error message should be displayed if the requested logical volume size cannot be created

[Could not create logical volume of that size](#)

[, and a size of 800 MiB](#) should be used instead .

If the volume group [research](#) does not exist, an error message should be displayed

[Volume group does not exist](#)

.

Do not mount logical volumes in any way

Problem solving method:

```
[greg@bastion ansible]$ vim lv.yml
---
- name: create and use logical volumes
  hosts: all
  tasks:
    - block:
        - name: one
          lvol:
            vg: research
            lv: data
            size: 1500
        - name: two
          filesystem:
            fstype: ext4
            dev: /dev/research/data
    rescue:
      - debug:
          msg: Could not create logical volume of that size
      - name: three
        lvol:
          vg: research
          lv: data
          size: 800
        when: ansible_lvm.vgs.research is defined
      - debug:
          msg: Volume group done not exist
        when: ansible_lvm.vgs.research is undefined
```

Generate hosts file

Download an initial template file from <http://materials/hosts.j2> to `/home/greg/ansible`

Complete the template so that it generates the following file: one line for each inventory host, in the same format as `/etc/hosts`

Create a playbook named `/home/greg/ansible/hosts.yml` that will use this template to generate the file `/etc/myhosts` on hosts in the dev host group .

After the playbook is running, the file `/etc/myhosts` on the hosts in the dev host group should contain one line for each managed host:

127.0.0.1 localhost localhost.localdomain localhost4 localhost4.localdomain4

::1 localhost localhost.localdomain localhost6 localhost6.localdomain6

172.25.250.9 workstation.lab.example.com workstation

172.25.250.10 servera.lab.example.com servera

172.25.250.11 serverb.lab.example.com serverb

172.25.250.12 serverc.lab.example.com serverc

172.25.250.13 serverd.lab.example.com serverd

Note: The order in which the inventory host names are displayed is not important.

Lao Liu's warm reminder: At present, there is no absolute good or bad for problem-solving methods A and B. It is easier to use A, but a small number of students report that points will be deducted. of. The B method is the standard answer of Red Hat's original factory, which is more reliable in theory, but the disadvantage is that the memory is complicated. Students are requested to make choices based on their own preparation progress and their own circumstances.

Solution A:

```
[greg@bastion ansible]$ wget http://materials/hosts.j2
[greg@bastion ansible]$ vim hosts.j2
127.0.0.1 localhost localhost.localdomain localhost4 localhost4.localdomain4
::1 localhost localhost.localdomain localhost6 localhost6.localdomain6
172.25.250.9 workstation.lab.example.com workstation
172.25.250.10 servera.lab.example.com servera
172.25.250.11 serverb.lab.example.com serverb
172.25.250.12 serverc.lab.example.com serverc
172.25.250.13 serverd.lab.example.com serverd
[greg@bastion ansible]$ vim hosts.yml
---
- name: generate hosts file
  hosts: dev
  tasks:
    - name: one
      template:
        src: /home/greg/ansible/hosts.j2
        dest: /etc/myhosts
```

Problem solving method:

```
[greg@bastion ansible]$ wget http://materials/hosts.j2
[greg@bastion ansible]$ vim hosts.j2
127.0.0.1 localhost localhost.localdomain localhost4 localhost4.localdomain4
::1 localhost localhost.localdomain localhost6 localhost6.localdomain
{% for host in groups['all'] %}
{{ hostvars[host]['ansible_facts']['default_ipv4']['address'] }} {{
hostvars[host]['ansible_facts']['fqdn'] }} {{ hostvars[host]['ansible_facts
']['hostname'] }}
{% end for %}
[greg@bastion ansible]$ vim hosts.yml
---
- name: generate hosts file
  hosts: dev
  tasks:
    - name: one
      template:
        src: /home/greg/ansible/hosts.j2
        dest: /etc/myhosts
```

Modify file content

Create a playbook named `/home/greg/ansible/issue.yml` as described below :

The playbook will run on all inventory hosts

The playbook will replace the contents of `/etc/issue` with a line of text shown below:

On hosts in the **dev** host group, this line of text reads: **Development**

On hosts in the **test** host group, this line of text reads: **Test**

On hosts in the **prod** host group, this line of text reads: **Production**

Problem solving method:

```
[greg@bastion ansible]$ vim issue.yml
---
- name: modify file content
  hosts: all
  tasks:
    - name: one
      copy:
        content: 'Development'
        dest: /etc/issue
      when: "inventory_hostname in groups.dev"
    - name: two
      copy:
        content: 'Test'
        dest: /etc/issue
      when: "inventory_hostname in groups.test"
    - name: three
      copy:
        content: 'Production'
        dest: /etc/issue
      when: "inventory_hostname in groups.prod"
```

😊 Create Web Content Directory

Create a playbook named **/home/greg/ansible/webcontent.yml** as described below :

The playbook runs on managed nodes in the **dev host group**

Create a directory **/webdev** that meets the following requirements :

The owner is the **webdev** group

With regular permissions: **owner=read+write+execute** , **group=read+write+execute** , **other=read+execute**

Has **special permissions** : set group ID

symlink **/var/www/html/webdev** to **/webdev**

Create the file **/webdev/index.html** with a one-line file like this: **Development**

Browsing this directory on a host in the dev host group (eg **http://172.25.250.9/webdev/**) will produce the following output:

Development

Problem solving method:

```
[greg@bastion ansible]$ vim webcontent.yml
```

```

---
- name: create web content directory
  hosts: dev
  tasks:
    - name: one
      file:
        path: /webdev
        state: directory
        group: webdev
        mode: '2775'
    - name: two
      file:
        src: /webdev
        dest: /var/www/html/webdev
        state: link
    - name: three
      copy:
        content: 'Development'
        dest: /webdev/index.html
        setype: httpd_sys_content_t

```

😊 Configure cron jobs (**new**)

Create a playbook called **/home/greg/ansible/cron.yml** ,

Configure a `cron` job that runs `every 2 minutes` and executes the following commands:

`logger "EX294 in progress"`, running as user `natasha`

Problem solving method:

```

[greg@bastion ansible]$ vim cron.yml
---
- name: configure cron job
  hosts: all
  tasks:
    - name: one
      user:
        name: natasha
        state: present
    - name: two
      cron:
        name: "logger info"
        minute: "*/2"
        user: natasha
        job: logger "EX294 in progress"

```

😊 Generate hardware report

Create a playbook named **/home/greg/ansible/hwreport.yml** that will generate an output file **/root/hwreport.txt** on all managed nodes with the following information :

Inventory hostname

Total memory size in MB

BIOS version

size of disk device **vda**

size of disk device **vdb**

Each line in the output file contains a key=value pair.

Your playbook should:

Download the file from <http://materials/hwreport.empty> and save it as **/root/hwreport.txt**

Change to **/root/hwreport.txt** with **the correct value**

If the hardware item does not exist, the associated value should be set to **NONE**

Problem solving method:

```
[greg@bastion ansible]$ vim hwreport.yml
---
- name: generate hardware report
  hosts: all
  vars:
    hw_all:
      - hw_name: HOST
        hw_cont: "{{ inventory_hostname | default('NONE',true) }}"
      - hw_name: MEMORY
        hw_cont: "{{ ansible_memtotal_mb | default('NONE',true) }}"
      - hw_name: BIOS
        hw_cont: "{{ ansible_bios_version | default('NONE',true) }}"
      - hw_name: DISK_SIZE_VDA
        hw_cont: "{{ ansible_devices.vda.size | default('NONE',true) }}"
      - hw_name: DISK_SIZE_VDB
        hw_cont: "{{ ansible_devices.vdb.size | default('NONE',true) }}"
  tasks:
    - name: one
      get_url:
        url: http://materials/hwreport.empty
        dest: /root/hwreport.txt
    - name: two
      lineinfile:
        path: /root/hwreport.txt
        regexp: "^{{ item.hw_name }}"
        line: "{{ item.hw_name }}={{ item.hw_cont }}"
      loop: "{{ hw_all }}"
```

Create password vault

Create an Ansible library to store user passwords as described below:

The library name is **/home/greg/ansible/locker.yml**

The library contains two variables with the following names:

pw_developer , the value is **Imadev**

pw_manager , the value is **Imamgr**

The password used to encrypt and decrypt the library is **whenyouwishuponastar**

The password is stored in the file **/home/greg/ansible/secret.txt**

Problem solving method:

```
[greg@bastion ansible]$ vim ansible.cfg
#Modify only one line
vault_password_file = /home/greg/ansible/secret.txt
[greg@bastion ansible]$ vim locker.yml
---
pw_developer: Imadev
pw_manager: Imamgr
[greg@bastion ansible]$ echo whenyouwishuponastar > secret.txt
[greg@bastion ansible]$ ansible-vault encrypt locker.yml
```

Create user account

Download the list of users to create from **http://materials/user_list.yml** and save it to **/home/greg/ansible**

The password library **/home/greg/ansible/locker.yml** created elsewhere is used for this exam . Create a playbook named **/home/greg/ansible/users.yml** to create user accounts as follows:

Users with the job description **developer** should:

Created on managed nodes in **dev** and **test host groups**

Assign password from **pw_developer** variable

is a member of the supplementary group **devops**

A user whose job description is **manager** should:

Created on a managed node in the **prod** host group

Assign password from **pw_manager** variable

is a member of the supplementary group **opsmgr**

The password is in **SHA512** hash format.

Your playbook should function properly for this exam using the vault secret file **/home/greg/ansible/secret.txt** created elsewhere.

Problem solving method:

```
[greg@bastion ansible]$ wget http://materials/user_list.yml
[greg@bastion ansible]$ vim users.yml
---
- name: create user account
  hosts: all
  vars_files:
    - /home/greg/ansible/locker.yml
    - /home/greg/ansible/user_list.yml
```

```

tasks:
  - name: one
    group:
      name: devops
      loop: "{{ users }}"
      when: item.job == 'developer' and (inventory_hostname in groups.dev or
inventory_hostname in groups.test)
  - name: two
    user:
      name: "{{ item.name }}"
      password: "{{ pw_developer | password_hash('sha512','mysecretsalt') }}"
    groups: devops
    loop: "{{ users }}"
    when: item.job == 'developer' and (inventory_hostname in groups.dev or
inventory_hostname in groups.test)
  - name: three
    group:
      name: opsmgr
      loop: "{{ users }}"
      when: item.job == 'manager' and inventory_hostname in groups.prod
  - name: four
    user:
      name: "{{ item.name }}"
      password: "{{ pw_manager | password_hash('sha512','mysecretsalt') }}"
    groups: opsmgr
    loop: "{{ users }}"
    when: item.job == 'manager' and inventory_hostname in groups.prod

```

🤪 Update keys for Ansible repositories

Update keys for existing Ansible repositories as described below:

Download the Ansible library from <http://materials/salaries.yml> to `/home/greg/ansible`

The current library password is `insecure8sure`

The new library password is `bbs2you9527`

The library remains encrypted with the new password

Problem solving method:

```

[greg@bastion ansible]$ wget http://materials/salaries.yml
[greg@bastion ansible]$ ansible-vault rekey --ask-vault-pass salaries.yml
Vault password: Paste current password
New Vault password: Paste new password
Confirm New Vault password: Paste new password
[greg@bastion ansible]$ ansible-vault view salaries.yml

```

🤪 Create partitions (hidden questions, have a chance to appear)

On the balancers host, divide a new partition, the device is `/dev/vdd`, number 1, size 1500m, format it as ext4, mount it to the `/newpart1` directory, if the space is not enough, divide it into 800m, if there is no vdd, an error is reported

Problem solving method:

```

---
- name: create partition
  hosts: balancers
  tasks:
    - block:
        - name: one
          file:
            path: /newpart1
            state: directory
        - name: two
          parted:
            device: /dev/vdd
            number: 1
            state: present
            part_end: 1500MiB
        - name: three
          filesystem:
            fstype: ext4
            dev: /dev/vdd1
        - name: four
          mount:
            path: /newpart1
            src: /dev/vdd1
            fstype: ext4
            state: mounted

    rescue:
      - debug:
          msg: Could not create partation of that size
      - name: five
        parted:
          device: /dev/vdd
          number: 1
          state: present
          part_end: 800MiB
        when: ansible_facts.devices.vdd is defined
      - name: six
        filesystem:
          fstype: ext4
          dev: /dev/vdd1
        when: ansible_facts.devices.vdd is defined
      - name: seven
        mount:
          path: /newpart1
          src: /dev/vdd1
          fstype: ext4
          state: mounted
        when: ansible_facts.devices.vdd is defined
      - debug:
          msg: Disk does not exist
        when: ansible_facts.devices.vdd is undefined

```

Install RHEL character (hidden question, there is a chance to appear)

Install the RHEL role, and use the SELinux role, which is required to run on all nodes, with SELinux set to enforcing mode.

```

[greg@bastion ansible]$ dnf install rhel-system-roles -y
[greg@bastion ansible]$ cp -rf /usr/share/ansible/roles/linux-system-roles.selinux ./
[greg@bastion ansible]$ vim selinux.yml
#After installation, you can change it or not. If you change it, you can leave the rest,
and delete the excess.

```

```
- hosts: all
  vars:
    selinux_policy: targeted
    selinux_state: enforcing
  roles:
    - role: roles/rhel-system-roles.selinux
  tasks:
    - name: apply SELinux role
      block:
        - include_roles:
            name: roles/rhel-system-roles.selinux
      rescue:
        - name: check
          fail:
            when: not selinux_reboot_required
        - name: reboot
          reboot:
        - name: changed
          include_role:
            name: roles/rhel-system-roles.selinux
```